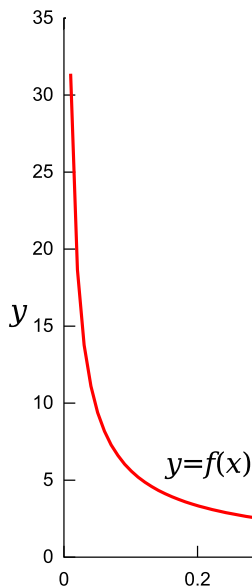


MPI profiling with Alinea MAP

Alexander Gaenko

February 5, 2016

Problem: Calculation of a definite integral.



The problem:

- We need to compute:

$$F(a, b) = \int_a^b f(x) dx$$

Where $f(x)$ is some (presumably “slow to compute”) function.

- We use this approach:

① Split $[a, b]$ by points $\{x_1, x_2, \dots, x_k\}$

②
$$F(a, b) = \sum_k f\left(\frac{x_k + x_{k+1}}{2}\right)(x_{k+1} - x_k)$$

- The integrand $f(x)$ and the integration routine are hidden inside a library.

Sequential program code

```
1  #include <stdio.h>
2  #include <math.h>
3
4  // Declare integrand() and integral() from 'mymath' library
5  #include "mylib/mymath.hpp"
6
7  int main(int argc, char** argv)
8  {
9      unsigned long int n;
10     if (argc!=2 || sscanf(argv[1], "%lu", &n)!=1) {
11         fprintf(stderr, "Usage: %s integration_steps\n\n", argv[0]);
12         return 1;
13     }
14
15     // Integration limits.
16     const double global_a=1E-5;
17     const double global_b=1;
18
19     // Perform integration
20     const double y=integral(integrand, n, global_a, global_b);
21
22     const double y_exact=4*(pow(global_b,0.25)-pow(global_a,0.25));
23     printf("Result=%lf Exact=%lf Difference=%lf\n", y, y_exact, y-y_exact);
24
25     return 0;
26 }
```

Problem size behavior: how to measure.

Do we even need to parallelize?


```
1 # Compilation:
2 $ gcc -O3 -o integral_seq.x integral_seq.cxx\
3     -L./mylib -lmymath
4 # Timed runs:
5 $ time -p ./integral_seq.x 1000000
6 Result=3.775045 Exact=3.775063 Difference=-0.000019
7 real 2.12
8 $ time -p ./integral_seq.x 2000000
9 Result=3.775058 Exact=3.775063 Difference=-0.000005
10 real 4.28
11 $ time -p ./integral_seq.x 8000000
12 Result=3.775062 Exact=3.775063 Difference=-0.000001
13 real 17.61
```

The time grows linearly with the problems size. Acceptable accuracy at 8M points. Can we speed it up?

Parallelization: domain decomposition.

Approach:

- Split $[a, b]$ into several domains;
- Compute integrals **independently**.



$$F(a, b) = \int_a^b f(x) dx$$

- 1 Assign a process to each domain $[x_k, x_{k+1}]$
- 2 Let each process compute $F(x_k, x_{k+1})$
- 3 $F(a, b) = \sum_k F(x_k, x_{k+1})$

“Embarassingly parallel” problem,
high speedup is expected.

A sketch of the parallel code

```
1  MPI_Init(&argc, &argv);
2
3  int rank, nprocs;
4  MPI_Comm_rank(MPI_COMM_WORLD, &rank);
5  MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
6
7  // ...Get total number of steps, broadcast it...
8  // ...
9
10 // Each rank figures out its integration limits and number of steps
11 unsigned long my_stepbase, my_nsteps;
12 get_steps(nsteps_all, nprocs, rank, &my_stepbase, &my_nsteps);
13
14 const double per_step=(global_b-global_a)/nsteps_all;
15 const double x1=global_a + my_stepbase*per_step;
16 const double x2=x1 + my_nsteps*per_step;
17
18 // Compute my own part of the integral
19 double my_y=integral(integrand, my_nsteps, x1, x2);
20
21 // Sum all numbers on master
22 double y=0;
23 MPI_Reduce(&my_y, &y, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
24
25 // ... print results ...
26 // ...
27
28 MPI_Barrier(MPI_COMM_WORLD);
29 // Here we could start another computation.
30 MPI_Finalize();
```

Parallel performance: how to measure.

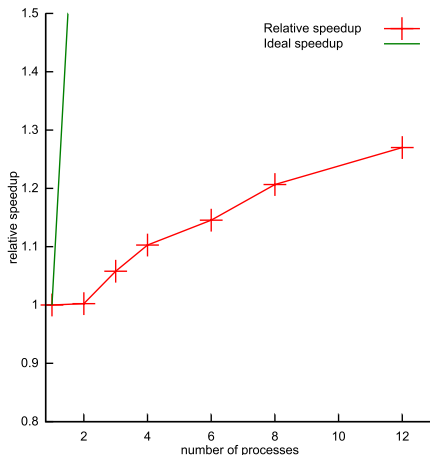
Now let's see how much we achieved...

- **Strong scaling**: as we add processes, how do we fare?
- **Weak scaling**: as we add *both* processes and work?

```
1 $ mpicc -O3 -o integral_par.x integral_par.cxx \  
2     -L./mylib -lmymath  
3 $ time -p mpirun -np 1 ./integral_par.x 8000000  
4 Result=3.775062 Exact=3.775063 Difference=-0.000001  
5 real 17.23  
6 user 17.08  
7 sys 0.02  
8 $ time -p mpirun -np 2 ./integral_par.x 8000000  
9 Result=3.775062 Exact=3.775063 Difference=-0.000001  
10 real 17.24  
11 user 31.98  
12 sys 0.05
```

Parallel performance: results

Is there a performance problem?



- Relative speedup:
$$s(p) = \frac{(\text{time with 1 process})}{(\text{time with } p \text{ processes})}$$
- Ideal relative speedup:
$$s_{\text{ideal}}(p) = p$$
- Our speedup is
25% on 12 nodes!
- I'd call it "*dismal*".
We **do** have a problem!
- Why? How to figure it out?

How does performance analysis work?

How to collect data?

- **Instrumentation:**
 - Insert timers & counters in the code
 - Requires source or binary processing
- **Sampling:**
 - Interrupt & check the program at regular intervals
 - Introduces statistical error

What kind of data?

- **Profile:**
 - Summary information only
 - Relatively small file
- **Trace:**
 - Detailed recording during the run
 - Potentially huge file
 - Profile can be restored

Allinea MAP does **tracing** by **sampling**.

Prepare for profiling

To prepare for profiling/tracing, one needs to:

- Compile with full optimization
- Generate debugging symbols
- Link with system libs dynamically
 - Usually the default
 - Notable exception: Cray
- On Flux: load ddt module

```
1 $ mpicc -g -O3 -o integral_par.x \  
2     integral_par.cxx -L ./mylib -lmymath  
3 $ module add ddt
```

Running Map: simple way (demo)

- 1 Get interactive access to a compute node
- 2 Change to your working directory
- 3 Optionally, set *sampling interval*
- 4 Run as you would, prefixed by `map`

```
1 $ qsub -V -I -X -q flux -l qos=flux,nproc=12 \  
2     -l walltime=10:0:0 -A account_flux  
3 $ cd $PBS_O_WORKDIR  
4 $ export ALLINEA_SAMPLER_INTERVAL=5  
5 $ map mpirun -np 12 ./integral_par.x 10000
```

Caution:

- Too small interval: large overhead!
- Too large interval: not enough samples!
- *Allinea* recommends at least 1000 samples/process

Running Map: other options

What if you can not or would not run a GUI?

- Have slow or non-existing X connection to compute nodes.
- Do not want to wait for interactive session.

Use `-profile` option.

```
1 #PBS -V
2 #PBS -q flux -l qos=flux -A account_flux
3 #PBS -l nproc=12,walltime=10:0:0
4 cd $PBS_O_WORKDIR
5 export ALLINEA_SAMPLER_INTERVAL=5
6 map -profile mpirun -np 12 ./integral_par.x 8000000
```

This will create a `*.map` file. Then run from the login node:

```
1 $ map integral_par_even_12p_*.map
```

If you are submitting to a Flux queue...

① Run map from the login node:

```
$ map ./integral_par.x 8000000
```

② Set number of processes

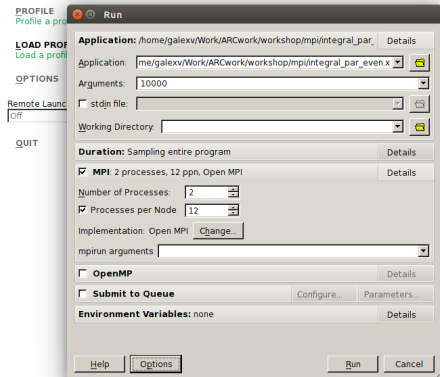
③ Check **Submit to queue**

④ Click **Configure...**

⑤ Load a proper *submission template file* (see next page)

⑥ Click **OK**

⑦ Click **Run**

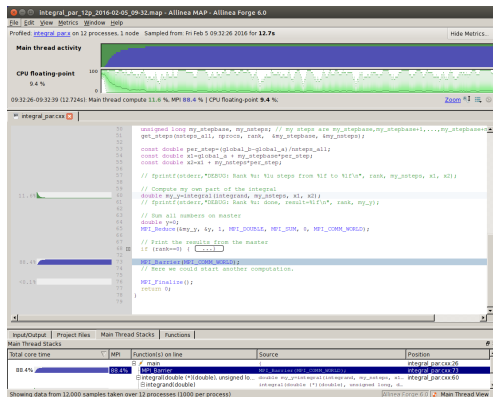


Submission template for Flux

```
1 #PBS -V
2 #PBS -l walltime=WALL_CLOCK_LIMIT_TAG
3 #PBS -l nodes=NUM_NODES_TAG:ppn=PROCS_PER_NODE_TAG
4 #PBS -q QUEUE_TAG -l qos=flux -A account_flux
5 #PBS -o PROGRAM_TAG-allinea.stdout
6 #PBS -e PROGRAM_TAG-allinea.stderr
7
8 cd $PBS_O_WORKDIR
9 AUTO_LAUNCH_TAG
```

Time for a live demo!

- Most of the time is spent in MPI;
- As run progresses, *even more* time is spent in MPI;
- Problem: some processes spend more time computing $f(x)$ then others!
- It's called “**Load Imbalance**”.



Possible solutions:

- Distribute work unevenly (but how?)
- Implement *Dynamic Load Balancing*.

Dynamic load balancing

Idea: If a process has nothing to do, make it to do something.

Manager-Workers approach:

Manager

- 1 Listen to all workers
- 2 Worker sent **READY** ?
 - send **GO** with a job chunk
- 3 Worker sent **DONE**?
 - add result to the sum
- 4 No more job chunks?
 - send **STOP** to the worker
- 5 No more workers?
 - we are done!
 - Otherwise, go to (1)

Worker


- 1 Send **READY** to the Manager
- 2 Listen to the Manager
- 3 Manager sent **GO** ?
 - Get job chunk
 - Do the calculation
 - Send **DONE** with result to the Manager
 - Go to (1)
- 4 Manager sent **STOP**?
 - exit.

Dynamic Load Balancing: large block size (2500)

Run with:

```
$ map mpirun -np 12 ./dyn_integral.x 8000000 2500
```


- CPU occupied till half-way
- Spikes in MPI use:
Manager receiving data
- Looks like the last worker was holding everyone
- Other workers:
worker starvation
(no work to do)



Dynamic Load Balancing: small block size (2)

Run with:

```
$ map mpirun -np 12 ./dyn_integral.x 8000000 2
```



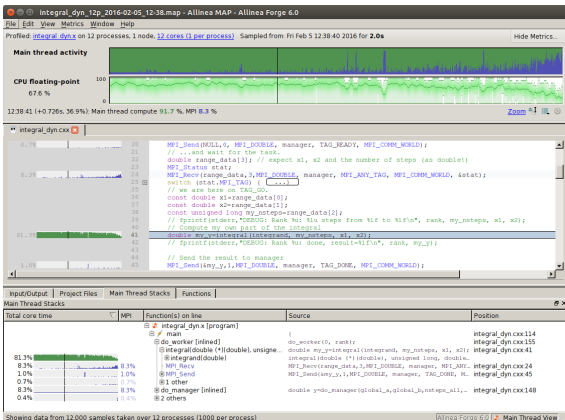
- Brief useful CPU work, then all time is in MPI
- Just moving data around:
 - Workers receiving data
 - 8% (1/12) of time: Manager sending data.
- Very low *computation/communication* ratio.

Dynamic Load Balancing: good block size (100)

Run with:


```
$ map mpirun -np 12 ./dyn_integral.x 8000000 100
```

- 80% time CPU is busy!
- 8% (1/12) of time: Manager work.
- Mostly MPI by the end of the run.
- OK computation/communication ratio.
- Still room for improvement!



Dynamic Load Balancing: Strong scaling graph.

How does it look now?



Conclusions:

- Block size does affect performance.
- Block size 100 grows up to node size (12).
- Too small block: MPI communication overhead.
- Too large block: workers starvation.

Concluding slide

Take-home message

- Once you made your program parallel, do simple scaling experiments.
- If scaling is bad, use profiling tools to understand why.
- *Allinea Map* is available for all Flux users.
- *Map* can analyze not only parallel, but single-node and OpenMP performance. (Can show something if time permits!)
- If you need any advise and/or help with your parallel programming, ARC provides [free consulting service](#)
 - Just send a mail to HPC support...
 - ...or directly to: "Alexander Gaenko" <galexv@umich.edu>

Thank you for your attention!